

#Data Preprocessing and Exploration

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.io import arff
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, roc_curve, confusion_matrix, classification_report
)
from keras.models import Model
from keras.layers import Input, Dense, Dropout
from keras.regularizers import l2
from keras.optimizers import Adam

# Load the .arff data
def load_arff_data(filename):
    data, meta = arff.loadarff(filename)
    df = pd.DataFrame(data)
    return df

# Load your training and test datasets
train_df = load_arff_data('ECG5000_TRAIN.arff')
test_df = load_arff_data('ECG5000_TEST.arff')

print(train_df.columns)
print('\n\n', train_df.describe())
print('\n\n', train_df.isna().sum())
print('\n\n', test_df.columns)
print('\n\n', test_df.describe())
print('\n\n', test_df.isna().sum())

Index(['att1', 'att2', 'att3', 'att4', 'att5', 'att6', 'att7', 'att8',
      'att9',
      'att10',
      ...,
      'att132', 'att133', 'att134', 'att135', 'att136', 'att137',
      'att138',
      'att139', 'att140', 'target'],
      dtype='object', length=141)
```

```
<bound method NDFrame.describe of
att4      att5      att6      att7 \
0   -0.112522 -2.827204 -3.773897 -4.349751 -4.376041 -3.474986 -
2.181408
1   -1.100878 -3.996840 -4.285843 -4.506579 -4.022377 -3.234368 -
```

```

1.566126
2   -0.567088 -2.593450 -3.874230 -4.584095 -4.187449 -3.151462 -
1.742940
3    0.490473 -1.914407 -3.616364 -4.318823 -4.268016 -3.881110 -
2.993280
4    0.800232 -0.874252 -2.384761 -3.973292 -4.338224 -3.802422 -
2.534510
..      ...      ...      ...      ...      ...      ...
...
495 -0.478577 -1.779959 -2.398159 -3.170112 -3.559732 -3.573956 -
2.989770
496 -1.325210 -2.480992 -2.965356 -3.342392 -3.176351 -2.891528 -
2.369679
497 -0.021964 -0.912434 -1.903353 -2.662829 -3.122156 -3.451490 -
3.392982
498  0.288011 -1.098020 -2.500250 -3.598599 -3.650608 -3.281587 -
2.231601
499 -1.133674 -2.702941 -3.120979 -3.558669 -3.312442 -2.607641 -
1.354939

      att8      att9      att10  ...      att132      att133
att134 \
0   -1.818286 -1.250522 -0.477492  ...   0.792168  0.933541  0.796958
1   -0.992258 -0.754680  0.042321  ...   0.538356  0.656881  0.787490
2   -1.490659 -1.183580 -0.394229  ...   0.886073  0.531452  0.311377
3   -1.671131 -1.333884 -0.965629  ...   0.350816  0.499111  0.600345
4   -1.783423 -1.594450 -0.753199  ...   1.148884  0.958434  1.059025
..      ...      ...      ...  ...      ...      ...      ...
495 -2.270605 -1.688277 -1.359872  ...   1.160885  1.456331  2.209421
496 -1.598750 -1.071751 -0.891843  ...  -0.172154 -0.864803 -1.549854
497 -2.929937 -2.256294 -1.690706  ...   1.339479  1.457995  2.128078
498 -1.250656 -1.072574 -0.434310  ...  -0.029242  0.071414  0.118161
499 -1.014740 -0.796023 -0.259599  ...  -3.206942 -2.941677 -2.557140

      att135      att136      att137      att138      att139      att140
target
0   0.578621  0.257740  0.228077  0.123431  0.925286  0.193137
b'1'
1   0.724046  0.555784  0.476333  0.773820  1.119621 -1.436250

```

```

b'1'
2 -0.021919 -0.713683 -0.532197  0.321097  0.904227 -0.421797
b'1'
3  0.842069  0.952074  0.990133  1.086798  1.403011 -0.383564
b'1'
4  1.371682  1.277392  0.960304  0.971020  1.614392  1.421456
b'1'
..      ...      ...      ...      ...      ...      ...      ..
.
495  2.507175  2.198534  1.705849  1.492642  1.561890  1.520161
b'4'
496 -2.460243 -3.366562 -3.466546 -2.718380 -1.855209 -1.539958
b'4'
497  2.630759  2.295748  1.764967  1.444280  1.432347  1.457028
b'4'
498 -0.071967 -0.171214  0.131211  0.049872  0.010915 -0.081534
b'5'
499 -1.487946 -1.118880 -0.737113 -0.110840  0.001858 -0.122639
b'5'

[500 rows x 141 columns]>

```

```

att1      0
att2      0
att3      0
att4      0
att5      0
..
att137    0
att138    0
att139    0
att140    0
target    0
Length: 141, dtype: int64

```

```

Index(['att1', 'att2', 'att3', 'att4', 'att5', 'att6', 'att7',
      'att8', 'att9',
      'att10',
      ...,
      'att132', 'att133', 'att134', 'att135', 'att136', 'att137',
      'att138',
      'att139', 'att140', 'target'],
      dtype='object', length=141)

```

```

<bound method NDFrame.describe of      att1      att2      att3
att4      att5      att6      att7  \
0      3.690844  0.711414 -2.114091 -4.141007 -4.574472 -3.431909 -

```

```

1.950791
1      -1.348132 -3.996038 -4.226750 -4.251187 -3.477953 -2.228422 -
1.808488
2      1.024295 -0.590314 -1.916949 -2.806989 -3.527905 -3.638675 -
2.779767
3      0.545657 -1.014383 -2.316698 -3.634040 -4.196857 -3.758093 -
3.194444
4      0.661133 -1.552471 -3.124641 -4.313351 -4.017042 -3.005993 -
1.832411

```

```

...      ...      ...      ...      ...      ...      ...
...
4495 -1.122969 -2.252925 -2.867628 -3.358605 -3.167849 -2.638360 -
1.664162
4496 -0.547705 -1.889545 -2.839779 -3.457912 -3.929149 -3.966026 -
3.492560
4497 -1.351779 -2.209006 -2.520225 -3.061475 -3.065141 -3.030739 -
2.622720
4498 -1.124432 -1.905039 -2.192707 -2.904320 -2.900722 -2.761252 -
2.569705
4499  0.728813  0.192597 -0.733884 -1.779456 -2.345908 -2.977565 -
3.380053

```

	att8	att9	att10	...	att132	att133	att134
\							
0	-1.107067	-0.632322	0.334577	...	0.022847	0.188937	0.480932
1	-1.534242	-0.779861	-0.397999	...	1.570938	1.591394	1.549193
2	-2.019031	-1.980754	-1.440680	...	0.443502	0.827582	1.237007
3	-2.221764	-1.588554	-1.202146	...	0.777530	1.119240	0.902984
4	-1.503886	-1.071705	-0.521316	...	1.280823	1.494315	1.618764
...
4495	-0.935655	-0.866953	-0.645363	...	-0.472419	-1.310147	-2.029521
4496	-2.695270	-1.849691	-1.374321	...	1.258419	1.907530	2.280888
4497	-2.044092	-1.295874	-0.733839	...	-1.512234	-2.076075	-2.586042
4498	-2.043893	-1.490538	-0.938473	...	-2.821782	-3.268355	-3.634981
4499	-3.417164	-3.030925	-2.313867	...	1.267275	1.678989	2.483389

	att135	att136	att137	att138	att139	att140
target						
0	0.629250	0.577291	0.665527	1.035997	1.492287	-1.905073

```

b'1'
1      1.193077  0.515134  0.126274  0.267532  1.071148 -1.164009
b'1'
2      1.235121  1.738103  1.800767  1.816301  1.473963  1.389767
b'1'
3      0.554098  0.497053  0.418116  0.703108  1.064602 -0.044853
b'1'
4      1.447449  1.238577  1.749692  1.986803  1.422756 -0.357784
b'1'
...      ...      ...      ...      ...      ...      ...
..
4495 -3.221294 -4.176790 -4.009720 -2.874136 -2.008369 -1.808334
b'4'
4496  1.895242  1.437702  1.193433  1.261335  1.150449  0.804932
b'2'
4497 -3.322799 -3.627311 -3.437038 -2.260023 -1.577823 -0.684531
b'2'
4498 -3.168765 -2.245878 -1.262260 -0.443307 -0.559769  0.108568
b'2'
4499  2.569073  2.122891  1.753963  1.538975  1.713781  1.309382
b'2'

[4500 rows x 141 columns]>

  att1      0
att2      0
att3      0
att4      0
att5      0
..
att137     0
att138     0
att139     0
att140     0
target     0
Length: 141, dtype: int64

```

Class Distribution

```

# Check unique classes and their counts in train and test
print("Class distribution in training set:")
print(train_df['target'].value_counts())

print("\nClass distribution in test set:")
print(test_df['target'].value_counts())

# Plot class distribution

```

```
plt.figure(figsize=(10, 4))

# Training set
plt.subplot(1, 2, 1)
sns.countplot(x='target', data=train_df, palette='Set2')
plt.title('Training Set Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')

# Test set
plt.subplot(1, 2, 2)
sns.countplot(x='target', data=test_df, palette='Set2')
plt.title('Test Set Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')

plt.tight_layout()
plt.show()
```

Class distribution in training set:

```
target
b'1'    292
b'2'    177
b'4'     19
b'3'     10
b'5'      2
Name: count, dtype: int64
```

Class distribution in test set:

```
target
b'1'    2627
b'2'    1590
b'4'     175
b'3'      86
b'5'      22
Name: count, dtype: int64
```

<ipython-input-3-52419079b109>:13: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='target', data=train_df, palette='Set2')
```

<ipython-input-3-52419079b109>:20: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='target', data=test_df, palette='Set2')
```



Test 1

```
# Decode the target labels from bytes to regular strings (e.g., b'1' → '1')
test_df['target'] = test_df['target'].apply(lambda x: x.decode('utf-8'))

# Extract the test features by dropping the target column
X_test_raw = test_df.drop(columns=['target']).values

# Create binary labels: 0 for normal (class '1'), 1 for all other classes (anomalies)
y_test = test_df['target'].apply(lambda x: 0 if x == '1' else 1).values

# Print the distribution of normal vs. anomalous samples in the test set
print("Final y_test distribution:", np.unique(y_test, return_counts=True))

Final y_test distribution: (array([0, 1]), array([2627, 1873]))
```

Autoencoder V1 (Base)

```
# Extract training features by removing the target column
X_train_raw = train_df.drop(columns=['target']).values

# Standardize the features: zero mean and unit variance
```

```

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train_raw)
X_test = scaler.transform(X_test_raw)

# Split a small portion of the training set for validation (10%)
X_train, X_val = train_test_split(X_train, test_size=0.1,
random_state=42)

def build_simple_autoencoder(input_dim):
    input_layer = Input(shape=(input_dim,))
    x = Dense(128, activation='relu')(input_layer)
    x = Dense(64, activation='relu')(x)
    x = Dense(32, activation='relu')(x)
    bottleneck = Dense(16, activation='relu')(x)
    x = Dense(32, activation='relu')(bottleneck)
    x = Dense(64, activation='relu')(x)
    x = Dense(128, activation='relu')(x)
    output_layer = Dense(input_dim, activation='linear')(x)
    autoencoder = Model(inputs=input_layer, outputs=output_layer)
    autoencoder.compile(optimizer=Adam(learning_rate=0.001),
loss='mse')
    return autoencoder

autoencoder = build_simple_autoencoder(X_train.shape[1])

history = autoencoder.fit(X_train, X_train, epochs=100, batch_size=32,
validation_data=(X_val, X_val), verbose=1)

Epoch 1/100
15/15 _____ 10s 62ms/step - loss: 0.9940 - val_loss:
1.0284
Epoch 2/100
15/15 _____ 1s 23ms/step - loss: 0.7516 - val_loss:
0.7580
Epoch 3/100
15/15 _____ 0s 28ms/step - loss: 0.6087 - val_loss:
0.5939
Epoch 4/100
15/15 _____ 1s 32ms/step - loss: 0.4574 - val_loss:
0.5169
Epoch 5/100
15/15 _____ 1s 26ms/step - loss: 0.3445 - val_loss:
0.4823
Epoch 6/100
15/15 _____ 1s 30ms/step - loss: 0.3702 - val_loss:
0.4607
Epoch 7/100
15/15 _____ 1s 33ms/step - loss: 0.3191 - val_loss:
0.4467
Epoch 8/100

```



```
15/15 ————— 0s 22ms/step - loss: 0.2963 - val_loss: 0.4233
Epoch 9/100
15/15 ————— 1s 20ms/step - loss: 0.2715 - val_loss: 0.4006
Epoch 10/100
15/15 ————— 1s 19ms/step - loss: 0.2482 - val_loss: 0.3751
Epoch 11/100
15/15 ————— 1s 15ms/step - loss: 0.2287 - val_loss: 0.3600
Epoch 12/100
15/15 ————— 0s 20ms/step - loss: 0.2191 - val_loss: 0.3480
Epoch 13/100
15/15 ————— 0s 25ms/step - loss: 0.1886 - val_loss: 0.3466
Epoch 14/100
15/15 ————— 0s 8ms/step - loss: 0.2088 - val_loss: 0.3429
Epoch 15/100
15/15 ————— 0s 8ms/step - loss: 0.1870 - val_loss: 0.3296
Epoch 16/100
15/15 ————— 0s 8ms/step - loss: 0.1789 - val_loss: 0.3232
Epoch 17/100
15/15 ————— 0s 8ms/step - loss: 0.1675 - val_loss: 0.3196
Epoch 18/100
15/15 ————— 0s 8ms/step - loss: 0.1796 - val_loss: 0.3238
Epoch 19/100
15/15 ————— 0s 9ms/step - loss: 0.1627 - val_loss: 0.3005
Epoch 20/100
15/15 ————— 0s 11ms/step - loss: 0.1491 - val_loss: 0.3084
Epoch 21/100
15/15 ————— 0s 13ms/step - loss: 0.1508 - val_loss: 0.3025
Epoch 22/100
15/15 ————— 0s 13ms/step - loss: 0.1562 - val_loss: 0.2984
Epoch 23/100
15/15 ————— 0s 14ms/step - loss: 0.1453 - val_loss: 0.3012
Epoch 24/100
15/15 ————— 0s 14ms/step - loss: 0.1444 - val_loss:
```

```
0.2883
Epoch 25/100
15/15 _____ 0s 15ms/step - loss: 0.1321 - val_loss:
0.2954
Epoch 26/100
15/15 _____ 0s 14ms/step - loss: 0.1332 - val_loss:
0.2911
Epoch 27/100
15/15 _____ 0s 16ms/step - loss: 0.1293 - val_loss:
0.2938
Epoch 28/100
15/15 _____ 0s 8ms/step - loss: 0.1362 - val_loss:
0.2813
Epoch 29/100
15/15 _____ 0s 9ms/step - loss: 0.1295 - val_loss:
0.2810
Epoch 30/100
15/15 _____ 0s 8ms/step - loss: 0.1271 - val_loss:
0.2777
Epoch 31/100
15/15 _____ 0s 9ms/step - loss: 0.1102 - val_loss:
0.2778
Epoch 32/100
15/15 _____ 0s 8ms/step - loss: 0.1274 - val_loss:
0.2692
Epoch 33/100
15/15 _____ 0s 8ms/step - loss: 0.1113 - val_loss:
0.2666
Epoch 34/100
15/15 _____ 0s 9ms/step - loss: 0.1087 - val_loss:
0.2871
Epoch 35/100
15/15 _____ 0s 8ms/step - loss: 0.1347 - val_loss:
0.2969
Epoch 36/100
15/15 _____ 0s 8ms/step - loss: 0.1276 - val_loss:
0.2707
Epoch 37/100
15/15 _____ 0s 9ms/step - loss: 0.1025 - val_loss:
0.2678
Epoch 38/100
15/15 _____ 0s 8ms/step - loss: 0.1051 - val_loss:
0.2598
Epoch 39/100
15/15 _____ 0s 8ms/step - loss: 0.1036 - val_loss:
0.2606
Epoch 40/100
15/15 _____ 0s 11ms/step - loss: 0.0993 - val_loss:
0.2656
```

```
Epoch 41/100
15/15 _____ 0s 8ms/step - loss: 0.1022 - val_loss:
0.2535
Epoch 42/100
15/15 _____ 0s 8ms/step - loss: 0.0946 - val_loss:
0.2568
Epoch 43/100
15/15 _____ 0s 8ms/step - loss: 0.1030 - val_loss:
0.2466
Epoch 44/100
15/15 _____ 0s 10ms/step - loss: 0.0928 - val_loss:
0.2530
Epoch 45/100
15/15 _____ 0s 9ms/step - loss: 0.0902 - val_loss:
0.2473
Epoch 46/100
15/15 _____ 0s 8ms/step - loss: 0.0957 - val_loss:
0.2470
Epoch 47/100
15/15 _____ 0s 8ms/step - loss: 0.0985 - val_loss:
0.2590
Epoch 48/100
15/15 _____ 0s 9ms/step - loss: 0.1104 - val_loss:
0.2735
Epoch 49/100
15/15 _____ 0s 8ms/step - loss: 0.1053 - val_loss:
0.2478
Epoch 50/100
15/15 _____ 0s 8ms/step - loss: 0.0921 - val_loss:
0.2397
Epoch 51/100
15/15 _____ 0s 9ms/step - loss: 0.0859 - val_loss:
0.2365
Epoch 52/100
15/15 _____ 0s 8ms/step - loss: 0.0921 - val_loss:
0.2408
Epoch 53/100
15/15 _____ 0s 8ms/step - loss: 0.0840 - val_loss:
0.2340
Epoch 54/100
15/15 _____ 0s 8ms/step - loss: 0.0797 - val_loss:
0.2446
Epoch 55/100
15/15 _____ 0s 8ms/step - loss: 0.0808 - val_loss:
0.2372
Epoch 56/100
15/15 _____ 0s 8ms/step - loss: 0.0773 - val_loss:
0.2395
Epoch 57/100
```


```
15/15 _____ 0s 9ms/step - loss: 0.0778 - val_loss: 0.2341
Epoch 58/100
15/15 _____ 0s 7ms/step - loss: 0.0775 - val_loss: 0.2361
Epoch 59/100
15/15 _____ 0s 8ms/step - loss: 0.0757 - val_loss: 0.2352
Epoch 60/100
15/15 _____ 0s 8ms/step - loss: 0.0725 - val_loss: 0.2303
Epoch 61/100
15/15 _____ 0s 8ms/step - loss: 0.0753 - val_loss: 0.2352
Epoch 62/100
15/15 _____ 0s 8ms/step - loss: 0.0730 - val_loss: 0.2334
Epoch 63/100
15/15 _____ 0s 11ms/step - loss: 0.0729 - val_loss: 0.2467
Epoch 64/100
15/15 _____ 0s 8ms/step - loss: 0.0937 - val_loss: 0.2472
Epoch 65/100
15/15 _____ 0s 8ms/step - loss: 0.1023 - val_loss: 0.2259
Epoch 66/100
15/15 _____ 0s 8ms/step - loss: 0.0732 - val_loss: 0.2199
Epoch 67/100
15/15 _____ 0s 8ms/step - loss: 0.0735 - val_loss: 0.2232
Epoch 68/100
15/15 _____ 0s 8ms/step - loss: 0.0719 - val_loss: 0.2208
Epoch 69/100
15/15 _____ 0s 9ms/step - loss: 0.0693 - val_loss: 0.2197
Epoch 70/100
15/15 _____ 0s 12ms/step - loss: 0.0681 - val_loss: 0.2168
Epoch 71/100
15/15 _____ 0s 8ms/step - loss: 0.0654 - val_loss: 0.2187
Epoch 72/100
15/15 _____ 0s 11ms/step - loss: 0.0746 - val_loss: 0.2150
Epoch 73/100
15/15 _____ 0s 8ms/step - loss: 0.0636 - val_loss:
```

```
0.2170
Epoch 74/100
15/15 _____ 0s 9ms/step - loss: 0.0635 - val_loss:
0.2172
Epoch 75/100
15/15 _____ 0s 8ms/step - loss: 0.0713 - val_loss:
0.2209
Epoch 76/100
15/15 _____ 0s 9ms/step - loss: 0.0692 - val_loss:
0.2198
Epoch 77/100
15/15 _____ 0s 9ms/step - loss: 0.0683 - val_loss:
0.2201
Epoch 78/100
15/15 _____ 0s 8ms/step - loss: 0.0651 - val_loss:
0.2188
Epoch 79/100
15/15 _____ 0s 8ms/step - loss: 0.0631 - val_loss:
0.2189
Epoch 80/100
15/15 _____ 0s 12ms/step - loss: 0.0647 - val_loss:
0.2205
Epoch 81/100
15/15 _____ 0s 17ms/step - loss: 0.0646 - val_loss:
0.2183
Epoch 82/100
15/15 _____ 0s 15ms/step - loss: 0.0661 - val_loss:
0.2197
Epoch 83/100
15/15 _____ 0s 14ms/step - loss: 0.0640 - val_loss:
0.2131
Epoch 84/100
15/15 _____ 0s 14ms/step - loss: 0.0600 - val_loss:
0.2155
Epoch 85/100
15/15 _____ 0s 13ms/step - loss: 0.0612 - val_loss:
0.2162
Epoch 86/100
15/15 _____ 0s 15ms/step - loss: 0.0589 - val_loss:
0.2110
Epoch 87/100
15/15 _____ 0s 15ms/step - loss: 0.0605 - val_loss:
0.2145
Epoch 88/100
15/15 _____ 0s 16ms/step - loss: 0.0561 - val_loss:
0.2129
Epoch 89/100
15/15 _____ 0s 9ms/step - loss: 0.0568 - val_loss:
0.2125
```

```
Epoch 90/100
15/15 _____ 0s 8ms/step - loss: 0.0573 - val_loss:
0.2158
Epoch 91/100
15/15 _____ 0s 8ms/step - loss: 0.0591 - val_loss:
0.2129
Epoch 92/100
15/15 _____ 0s 9ms/step - loss: 0.0571 - val_loss:
0.2151
Epoch 93/100
15/15 _____ 0s 11ms/step - loss: 0.0574 - val_loss:
0.2224
Epoch 94/100
15/15 _____ 0s 9ms/step - loss: 0.0572 - val_loss:
0.2159
Epoch 95/100
15/15 _____ 0s 8ms/step - loss: 0.0588 - val_loss:
0.2218
Epoch 96/100
15/15 _____ 0s 8ms/step - loss: 0.0593 - val_loss:
0.2172
Epoch 97/100
15/15 _____ 0s 9ms/step - loss: 0.0556 - val_loss:
0.2136
Epoch 98/100
15/15 _____ 0s 8ms/step - loss: 0.0558 - val_loss:
0.2159
Epoch 99/100
15/15 _____ 0s 8ms/step - loss: 0.0561 - val_loss:
0.2141
Epoch 100/100
15/15 _____ 0s 8ms/step - loss: 0.0585 - val_loss:
0.2117
```

Loss Plot

```
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Training vs Validation Loss - Model 1')
plt.xlabel('Epochs'); plt.ylabel('MSE'); plt.legend(); plt.grid(True)
plt.show()
```



Testing Evaluation

```
X_test_reconstructed = autoencoder.predict(X_test)
reconstruction_errors = np.mean(np.square(X_test -
X_test_reconstructed), axis=1)
X_train_reconstructed = autoencoder.predict(X_train)
train_errors = np.mean(np.square(X_train - X_train_reconstructed),
axis=1)
threshold = np.percentile(train_errors, 90)
print(f"Anomaly Threshold (90th percentile): {threshold:.4f}")
```

```
141/141 _____ 0s 2ms/step
15/15 _____ 0s 3ms/step
Anomaly Threshold (90th percentile): 0.0898
```

```
y_pred = (reconstruction_errors > threshold).astype(int)
print("\n--- Evaluation Report (Model 1) ---")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f"Precision: {precision_score(y_test, y_pred):.4f}")
print(f"Recall: {recall_score(y_test, y_pred):.4f}")
print(f"F1: {f1_score(y_test, y_pred):.4f}")
```

```
print(f"ROC-AUC: {roc_auc_score(y_test, reconstruction_errors):.4f}")
print(classification_report(y_test, y_pred))
```

```
--- Evaluation Report (Model 1) ---
```

```
Accuracy: 0.5189
```

```
Precision: 0.3748
```

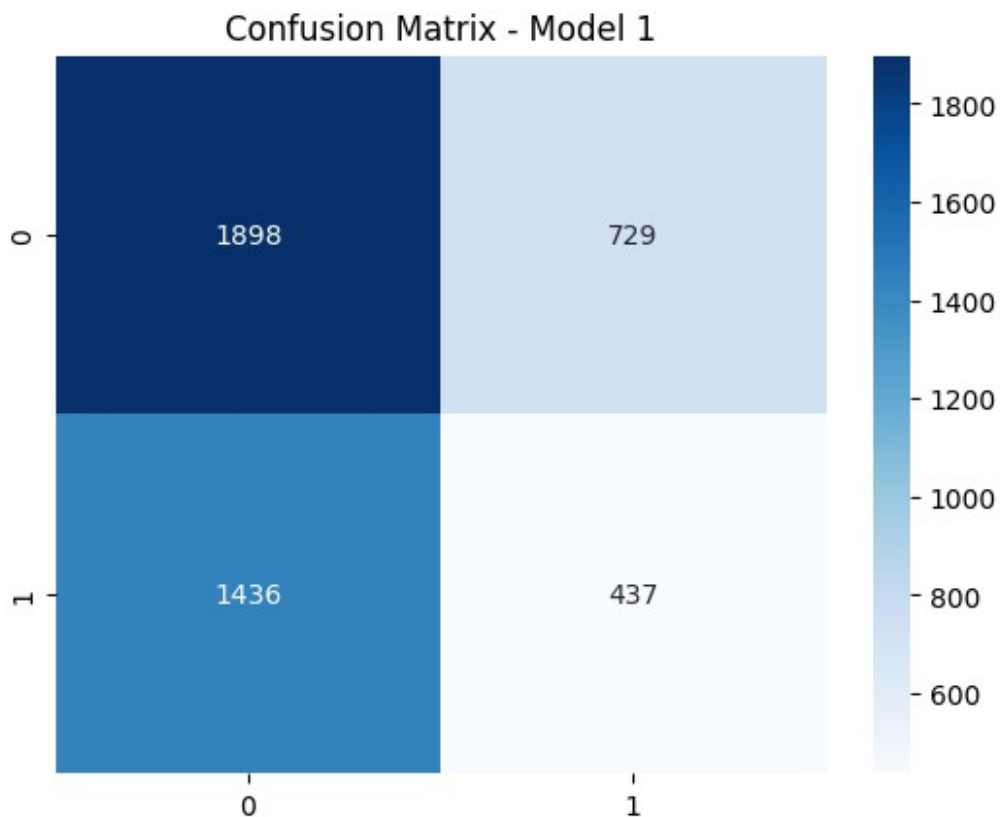
```
Recall: 0.2333
```

```
F1: 0.2876
```

```
ROC-AUC: 0.4230
```

	precision	recall	f1-score	support
0	0.57	0.72	0.64	2627
1	0.37	0.23	0.29	1873
accuracy			0.52	4500
macro avg	0.47	0.48	0.46	4500
weighted avg	0.49	0.52	0.49	4500

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True,
cmap='Blues', fmt='d')
plt.title("Confusion Matrix - Model 1")
plt.show()
```



Autoencoder V2 (Normal Only)

```
train_df['target'] = train_df['target'].apply(lambda x: x.decode('utf-8'))
train_normal_df = train_df[train_df['target'] == '1']
X_train_raw = train_normal_df.drop(columns=['target']).values
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train_raw)
X_test = scaler.transform(X_test_raw)
X_train, X_val = train_test_split(X_train, test_size=0.1,
random_state=42)
```

```
autoencoder = build_simple_autoencoder(X_train.shape[1])
```

```
history = autoencoder.fit(X_train, X_train, epochs=100, batch_size=32,
validation_data=(X_val, X_val), verbose=1)
```

Epoch 1/100

9/9 _____ 3s 37ms/step - loss: 0.9442 - val_loss: 0.8206

Epoch 2/100

9/9 _____ 0s 12ms/step - loss: 0.9301 - val_loss: 0.7340

Epoch 3/100

9/9 _____ 0s 12ms/step - loss: 0.8918 - val_loss: 0.6326

Epoch 4/100

9/9 _____ 0s 11ms/step - loss: 0.7002 - val_loss: 0.5255

Epoch 5/100

9/9 _____ 0s 12ms/step - loss: 0.5663 - val_loss: 0.4454

Epoch 6/100

9/9 _____ 0s 13ms/step - loss: 0.4245 - val_loss: 0.3768

Epoch 7/100

9/9 _____ 0s 16ms/step - loss: 0.3717 - val_loss: 0.3225

Epoch 8/100

9/9 _____ 0s 12ms/step - loss: 0.3344 - val_loss: 0.2999

Epoch 9/100

9/9 _____ 0s 11ms/step - loss: 0.3136 - val_loss: 0.2960

Epoch 10/100

9/9 _____ 0s 12ms/step - loss: 0.3160 - val_loss: 0.2706

Epoch 11/100

9/9 _____ 0s 12ms/step - loss: 0.2695 - val_loss: 0.2650

```
Epoch 12/100
9/9 _____ 0s 17ms/step - loss: 0.2687 - val_loss:
0.2535
Epoch 13/100
9/9 _____ 0s 17ms/step - loss: 0.2658 - val_loss:
0.2562
Epoch 14/100
9/9 _____ 0s 12ms/step - loss: 0.2631 - val_loss:
0.2465
Epoch 15/100
9/9 _____ 0s 12ms/step - loss: 0.2634 - val_loss:
0.2483
Epoch 16/100
9/9 _____ 0s 12ms/step - loss: 0.2427 - val_loss:
0.2390
Epoch 17/100
9/9 _____ 0s 12ms/step - loss: 0.2380 - val_loss:
0.2345
Epoch 18/100
9/9 _____ 0s 12ms/step - loss: 0.2291 - val_loss:
0.2282
Epoch 19/100
9/9 _____ 0s 19ms/step - loss: 0.2298 - val_loss:
0.2228
Epoch 20/100
9/9 _____ 0s 24ms/step - loss: 0.2208 - val_loss:
0.2215
Epoch 21/100
9/9 _____ 0s 17ms/step - loss: 0.2081 - val_loss:
0.2148
Epoch 22/100
9/9 _____ 0s 21ms/step - loss: 0.2090 - val_loss:
0.2114
Epoch 23/100
9/9 _____ 0s 26ms/step - loss: 0.1976 - val_loss:
0.2061
Epoch 24/100
9/9 _____ 0s 20ms/step - loss: 0.1964 - val_loss:
0.2008
Epoch 25/100
9/9 _____ 0s 24ms/step - loss: 0.1868 - val_loss:
0.1957
Epoch 26/100
9/9 _____ 0s 22ms/step - loss: 0.1766 - val_loss:
0.1956
Epoch 27/100
9/9 _____ 0s 22ms/step - loss: 0.1784 - val_loss:
0.1892
Epoch 28/100
```

```
9/9 _____ 0s 11ms/step - loss: 0.1753 - val_loss: 0.1861
Epoch 29/100
9/9 _____ 0s 12ms/step - loss: 0.1847 - val_loss: 0.1900
Epoch 30/100
9/9 _____ 0s 11ms/step - loss: 0.1710 - val_loss: 0.1852
Epoch 31/100
9/9 _____ 0s 17ms/step - loss: 0.1591 - val_loss: 0.1820
Epoch 32/100
9/9 _____ 0s 12ms/step - loss: 0.1478 - val_loss: 0.1794
Epoch 33/100
9/9 _____ 0s 17ms/step - loss: 0.1660 - val_loss: 0.1796
Epoch 34/100
9/9 _____ 0s 12ms/step - loss: 0.1483 - val_loss: 0.1762
Epoch 35/100
9/9 _____ 0s 12ms/step - loss: 0.1539 - val_loss: 0.1755
Epoch 36/100
9/9 _____ 0s 11ms/step - loss: 0.1550 - val_loss: 0.1749
Epoch 37/100
9/9 _____ 0s 12ms/step - loss: 0.1439 - val_loss: 0.1741
Epoch 38/100
9/9 _____ 0s 11ms/step - loss: 0.1453 - val_loss: 0.1741
Epoch 39/100
9/9 _____ 0s 17ms/step - loss: 0.1521 - val_loss: 0.1744
Epoch 40/100
9/9 _____ 0s 12ms/step - loss: 0.1508 - val_loss: 0.1705
Epoch 41/100
9/9 _____ 0s 11ms/step - loss: 0.1435 - val_loss: 0.1731
Epoch 42/100
9/9 _____ 0s 17ms/step - loss: 0.1384 - val_loss: 0.1702
Epoch 43/100
9/9 _____ 0s 12ms/step - loss: 0.1321 - val_loss: 0.1690
Epoch 44/100
9/9 _____ 0s 13ms/step - loss: 0.1377 - val_loss:
```

```
0.1679
Epoch 45/100
9/9 _____ 0s 12ms/step - loss: 0.1416 - val_loss:
0.1662
Epoch 46/100
9/9 _____ 0s 11ms/step - loss: 0.1284 - val_loss:
0.1671
Epoch 47/100
9/9 _____ 0s 11ms/step - loss: 0.1286 - val_loss:
0.1673
Epoch 48/100
9/9 _____ 0s 12ms/step - loss: 0.1320 - val_loss:
0.1627
Epoch 49/100
9/9 _____ 0s 11ms/step - loss: 0.1313 - val_loss:
0.1690
Epoch 50/100
9/9 _____ 0s 12ms/step - loss: 0.1275 - val_loss:
0.1641
Epoch 51/100
9/9 _____ 0s 19ms/step - loss: 0.1196 - val_loss:
0.1630
Epoch 52/100
9/9 _____ 0s 12ms/step - loss: 0.1333 - val_loss:
0.1661
Epoch 53/100
9/9 _____ 0s 13ms/step - loss: 0.1209 - val_loss:
0.1642
Epoch 54/100
9/9 _____ 0s 12ms/step - loss: 0.1195 - val_loss:
0.1677
Epoch 55/100
9/9 _____ 0s 12ms/step - loss: 0.1114 - val_loss:
0.1611
Epoch 56/100
9/9 _____ 0s 12ms/step - loss: 0.1206 - val_loss:
0.1623
Epoch 57/100
9/9 _____ 0s 17ms/step - loss: 0.1171 - val_loss:
0.1613
Epoch 58/100
9/9 _____ 0s 12ms/step - loss: 0.1180 - val_loss:
0.1648
Epoch 59/100
9/9 _____ 0s 12ms/step - loss: 0.1155 - val_loss:
0.1609
Epoch 60/100
9/9 _____ 0s 14ms/step - loss: 0.1246 - val_loss:
0.1625
```

Epoch 61/100
9/9 _____ 0s 12ms/step - loss: 0.1145 - val_loss: 0.1623
Epoch 62/100
9/9 _____ 0s 12ms/step - loss: 0.1311 - val_loss: 0.1657
Epoch 63/100
9/9 _____ 0s 17ms/step - loss: 0.1195 - val_loss: 0.1630
Epoch 64/100
9/9 _____ 0s 17ms/step - loss: 0.1125 - val_loss: 0.1538
Epoch 65/100
9/9 _____ 0s 12ms/step - loss: 0.1125 - val_loss: 0.1589
Epoch 66/100
9/9 _____ 0s 13ms/step - loss: 0.1024 - val_loss: 0.1542
Epoch 67/100
9/9 _____ 0s 13ms/step - loss: 0.1085 - val_loss: 0.1530
Epoch 68/100
9/9 _____ 0s 11ms/step - loss: 0.0995 - val_loss: 0.1493
Epoch 69/100
9/9 _____ 0s 12ms/step - loss: 0.0986 - val_loss: 0.1465
Epoch 70/100
9/9 _____ 0s 16ms/step - loss: 0.0928 - val_loss: 0.1514
Epoch 71/100
9/9 _____ 0s 16ms/step - loss: 0.1031 - val_loss: 0.1494
Epoch 72/100
9/9 _____ 0s 13ms/step - loss: 0.0992 - val_loss: 0.1499
Epoch 73/100
9/9 _____ 0s 12ms/step - loss: 0.0934 - val_loss: 0.1451
Epoch 74/100
9/9 _____ 0s 12ms/step - loss: 0.0968 - val_loss: 0.1490
Epoch 75/100
9/9 _____ 0s 12ms/step - loss: 0.0891 - val_loss: 0.1448
Epoch 76/100
9/9 _____ 0s 12ms/step - loss: 0.0902 - val_loss: 0.1453
Epoch 77/100

```
9/9 _____ 0s 13ms/step - loss: 0.0883 - val_loss:
0.1431
Epoch 78/100
9/9 _____ 0s 12ms/step - loss: 0.0943 - val_loss:
0.1453
Epoch 79/100
9/9 _____ 0s 11ms/step - loss: 0.0858 - val_loss:
0.1407
Epoch 80/100
9/9 _____ 0s 12ms/step - loss: 0.0872 - val_loss:
0.1416
Epoch 81/100
9/9 _____ 0s 12ms/step - loss: 0.0854 - val_loss:
0.1432
Epoch 82/100
9/9 _____ 0s 12ms/step - loss: 0.0873 - val_loss:
0.1417
Epoch 83/100
9/9 _____ 0s 12ms/step - loss: 0.0881 - val_loss:
0.1421
Epoch 84/100
9/9 _____ 0s 12ms/step - loss: 0.0844 - val_loss:
0.1436
Epoch 85/100
9/9 _____ 0s 13ms/step - loss: 0.0860 - val_loss:
0.1434
Epoch 86/100
9/9 _____ 0s 12ms/step - loss: 0.0824 - val_loss:
0.1405
Epoch 87/100
9/9 _____ 0s 12ms/step - loss: 0.0833 - val_loss:
0.1400
Epoch 88/100
9/9 _____ 0s 13ms/step - loss: 0.0858 - val_loss:
0.1403
Epoch 89/100
9/9 _____ 0s 12ms/step - loss: 0.0840 - val_loss:
0.1381
Epoch 90/100
9/9 _____ 0s 12ms/step - loss: 0.0853 - val_loss:
0.1402
Epoch 91/100
9/9 _____ 0s 12ms/step - loss: 0.0827 - val_loss:
0.1352
Epoch 92/100
9/9 _____ 0s 13ms/step - loss: 0.0845 - val_loss:
0.1365
Epoch 93/100
9/9 _____ 0s 12ms/step - loss: 0.0788 - val_loss:
```

```

0.1374
Epoch 94/100
9/9 _____ 0s 11ms/step - loss: 0.0760 - val_loss:
0.1365
Epoch 95/100
9/9 _____ 0s 17ms/step - loss: 0.0758 - val_loss:
0.1351
Epoch 96/100
9/9 _____ 0s 21ms/step - loss: 0.0842 - val_loss:
0.1357
Epoch 97/100
9/9 _____ 0s 21ms/step - loss: 0.0775 - val_loss:
0.1390
Epoch 98/100
9/9 _____ 0s 20ms/step - loss: 0.0790 - val_loss:
0.1354
Epoch 99/100
9/9 _____ 0s 19ms/step - loss: 0.0747 - val_loss:
0.1357
Epoch 100/100
9/9 _____ 0s 22ms/step - loss: 0.0770 - val_loss:
0.1367

```

Test Evaluation (V2)

```

X_test_reconstructed = autoencoder.predict(X_test)
reconstruction_errors = np.mean(np.square(X_test -
X_test_reconstructed), axis=1)
X_train_reconstructed = autoencoder.predict(X_train)
train_errors = np.mean(np.square(X_train - X_train_reconstructed),
axis=1)
threshold = np.percentile(train_errors, 90)
print(f"Anomaly Threshold (90th percentile): {threshold:.4f}")

141/141 _____ 1s 3ms/step
9/9 _____ 0s 4ms/step
Anomaly Threshold (90th percentile): 0.1178

y_pred = (reconstruction_errors > threshold).astype(int)
print("\n--- Evaluation Report (Model 2 - Normal Only) ---")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f"Precision: {precision_score(y_test, y_pred):.4f}")
print(f"Recall: {recall_score(y_test, y_pred):.4f}")
print(f"F1: {f1_score(y_test, y_pred):.4f}")
print(f"ROC-AUC: {roc_auc_score(y_test, reconstruction_errors):.4f}")
print(classification_report(y_test, y_pred))

```

--- Evaluation Report (Model 2 - Normal Only) ---

Accuracy: 0.7969

Precision: 0.6723


Recall: 0.9989

F1: 0.8037

ROC-AUC: 0.9857

	precision	recall	f1-score	support
0	1.00	0.65	0.79	2627
1	0.67	1.00	0.80	1873
accuracy			0.80	4500
macro avg	0.84	0.83	0.80	4500
weighted avg	0.86	0.80	0.80	4500

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True,  
cmap='Blues', fmt='d')  
plt.title("Confusion Matrix - Model 2")  
plt.show()
```



Autoencoder V3 (Normal + L2 + Dropout)

```
def build_regularized_autoencoder(input_dim):
    input_layer = Input(shape=(input_dim,))
    x = Dense(128, activation='relu', kernel_regularizer=l2(0.001))
    (input_layer)
    x = Dropout(0.2)(x)
    x = Dense(64, activation='relu', kernel_regularizer=l2(0.001))(x)
    x = Dropout(0.2)(x)
    x = Dense(32, activation='relu')(x)
    bottleneck = Dense(16, activation='relu')(x)
    x = Dense(32, activation='relu')(bottleneck)
    x = Dense(64, activation='relu')(x)
    x = Dense(128, activation='relu')(x)
    output_layer = Dense(input_dim, activation='linear')(x)
    model = Model(inputs=input_layer, outputs=output_layer)
    model.compile(optimizer=Adam(learning_rate=0.001), loss='mse')
    return model
```

```
autoencoder = build_regularized_autoencoder(X_train.shape[1])
```

```
history = autoencoder.fit(X_train, X_train, epochs=100, batch_size=32,
validation_data=(X_val, X_val), verbose=1)
```

Epoch 1/100

9/9 _____ 3s 43ms/step - loss: 1.1496 - val_loss: 1.0417

Epoch 2/100

9/9 _____ 0s 13ms/step - loss: 1.1396 - val_loss: 1.0055

Epoch 3/100

9/9 _____ 0s 16ms/step - loss: 0.9782 - val_loss: 0.9322

Epoch 4/100

9/9 _____ 0s 13ms/step - loss: 0.9161 - val_loss: 0.8641

Epoch 5/100

9/9 _____ 0s 12ms/step - loss: 0.8680 - val_loss: 0.8090

Epoch 6/100

9/9 _____ 0s 13ms/step - loss: 0.7799 - val_loss: 0.7012

Epoch 7/100

9/9 _____ 0s 14ms/step - loss: 0.6783 - val_loss: 0.6052

Epoch 8/100

9/9 _____ 0s 17ms/step - loss: 0.5900 - val_loss: 0.5086

Epoch 9/100

9/9 _____ 0s 12ms/step - loss: 0.6086 - val_loss:

```
0.4670
Epoch 10/100
9/9 _____ 0s 13ms/step - loss: 0.4884 - val_loss:
0.4369
Epoch 11/100
9/9 _____ 0s 13ms/step - loss: 0.5079 - val_loss:
0.4186
Epoch 12/100
9/9 _____ 0s 12ms/step - loss: 0.4714 - val_loss:
0.4043
Epoch 13/100
9/9 _____ 0s 13ms/step - loss: 0.4544 - val_loss:
0.3931
Epoch 14/100
9/9 _____ 0s 14ms/step - loss: 0.4646 - val_loss:
0.3887
Epoch 15/100
9/9 _____ 0s 17ms/step - loss: 0.4314 - val_loss:
0.3775
Epoch 16/100
9/9 _____ 0s 12ms/step - loss: 0.4887 - val_loss:
0.3723
Epoch 17/100
9/9 _____ 0s 13ms/step - loss: 0.4711 - val_loss:
0.3674
Epoch 18/100
9/9 _____ 0s 12ms/step - loss: 0.4199 - val_loss:
0.3569
Epoch 19/100
9/9 _____ 0s 12ms/step - loss: 0.4114 - val_loss:
0.3458
Epoch 20/100
9/9 _____ 0s 12ms/step - loss: 0.4094 - val_loss:
0.3428
Epoch 21/100
9/9 _____ 0s 17ms/step - loss: 0.3951 - val_loss:
0.3491
Epoch 22/100
9/9 _____ 0s 12ms/step - loss: 0.4256 - val_loss:
0.3436
Epoch 23/100
9/9 _____ 0s 12ms/step - loss: 0.4129 - val_loss:
0.3395
Epoch 24/100
9/9 _____ 0s 12ms/step - loss: 0.4272 - val_loss:
0.3287
Epoch 25/100
9/9 _____ 0s 12ms/step - loss: 0.3899 - val_loss:
0.3313
```

```
Epoch 26/100
9/9 _____ 0s 17ms/step - loss: 0.3822 - val_loss:
0.3344
Epoch 27/100
9/9 _____ 0s 12ms/step - loss: 0.3928 - val_loss:
0.3229
Epoch 28/100
9/9 _____ 0s 12ms/step - loss: 0.3635 - val_loss:
0.3191
Epoch 29/100
9/9 _____ 0s 17ms/step - loss: 0.3819 - val_loss:
0.3181
Epoch 30/100
9/9 _____ 0s 12ms/step - loss: 0.3607 - val_loss:
0.3202
Epoch 31/100
9/9 _____ 0s 12ms/step - loss: 0.3585 - val_loss:
0.3237
Epoch 32/100
9/9 _____ 0s 17ms/step - loss: 0.3686 - val_loss:
0.3187
Epoch 33/100
9/9 _____ 0s 13ms/step - loss: 0.3411 - val_loss:
0.3093
Epoch 34/100
9/9 _____ 0s 12ms/step - loss: 0.3510 - val_loss:
0.3019
Epoch 35/100
9/9 _____ 0s 12ms/step - loss: 0.3390 - val_loss:
0.3131
Epoch 36/100
9/9 _____ 0s 12ms/step - loss: 0.3424 - val_loss:
0.3026
Epoch 37/100
9/9 _____ 0s 17ms/step - loss: 0.3629 - val_loss:
0.2983
Epoch 38/100
9/9 _____ 0s 23ms/step - loss: 0.3157 - val_loss:
0.3049
Epoch 39/100
9/9 _____ 0s 21ms/step - loss: 0.3185 - val_loss:
0.3087
Epoch 40/100
9/9 _____ 0s 22ms/step - loss: 0.3144 - val_loss:
0.3084
Epoch 41/100
9/9 _____ 0s 21ms/step - loss: 0.3467 - val_loss:
0.3049
Epoch 42/100
```

```
9/9 _____ 0s 22ms/step - loss: 0.3108 - val_loss: 0.3000
Epoch 43/100
9/9 _____ 0s 24ms/step - loss: 0.3148 - val_loss: 0.3024
Epoch 44/100
9/9 _____ 0s 24ms/step - loss: 0.3204 - val_loss: 0.2968
Epoch 45/100
9/9 _____ 0s 23ms/step - loss: 0.3370 - val_loss: 0.2969
Epoch 46/100
9/9 _____ 0s 15ms/step - loss: 0.3054 - val_loss: 0.2940
Epoch 47/100
9/9 _____ 0s 12ms/step - loss: 0.3216 - val_loss: 0.2942
Epoch 48/100
9/9 _____ 0s 12ms/step - loss: 0.3219 - val_loss: 0.2839
Epoch 49/100
9/9 _____ 0s 12ms/step - loss: 0.3109 - val_loss: 0.2913
Epoch 50/100
9/9 _____ 0s 18ms/step - loss: 0.3080 - val_loss: 0.2828
Epoch 51/100
9/9 _____ 0s 14ms/step - loss: 0.3147 - val_loss: 0.2939
Epoch 52/100
9/9 _____ 0s 12ms/step - loss: 0.3006 - val_loss: 0.2861
Epoch 53/100
9/9 _____ 0s 12ms/step - loss: 0.2814 - val_loss: 0.2843
Epoch 54/100
9/9 _____ 0s 12ms/step - loss: 0.3034 - val_loss: 0.2883
Epoch 55/100
9/9 _____ 0s 12ms/step - loss: 0.2929 - val_loss: 0.2779
Epoch 56/100
9/9 _____ 0s 11ms/step - loss: 0.2833 - val_loss: 0.2857
Epoch 57/100
9/9 _____ 0s 12ms/step - loss: 0.3014 - val_loss: 0.2786
Epoch 58/100
9/9 _____ 0s 14ms/step - loss: 0.2822 - val_loss:
```

```
0.2767
Epoch 59/100
9/9 _____ 0s 18ms/step - loss: 0.2773 - val_loss:
0.2863
Epoch 60/100
9/9 _____ 0s 12ms/step - loss: 0.2924 - val_loss:
0.2787
Epoch 61/100
9/9 _____ 0s 11ms/step - loss: 0.2852 - val_loss:
0.2752
Epoch 62/100
9/9 _____ 0s 12ms/step - loss: 0.2702 - val_loss:
0.2692
Epoch 63/100
9/9 _____ 0s 13ms/step - loss: 0.2592 - val_loss:
0.2711
Epoch 64/100
9/9 _____ 0s 13ms/step - loss: 0.2670 - val_loss:
0.2664
Epoch 65/100
9/9 _____ 0s 12ms/step - loss: 0.2743 - val_loss:
0.2671
Epoch 66/100
9/9 _____ 0s 13ms/step - loss: 0.2702 - val_loss:
0.2661
Epoch 67/100
9/9 _____ 0s 12ms/step - loss: 0.2716 - val_loss:
0.2674
Epoch 68/100
9/9 _____ 0s 11ms/step - loss: 0.2737 - val_loss:
0.2745
Epoch 69/100
9/9 _____ 0s 12ms/step - loss: 0.2967 - val_loss:
0.2793
Epoch 70/100
9/9 _____ 0s 13ms/step - loss: 0.2749 - val_loss:
0.2639
Epoch 71/100
9/9 _____ 0s 13ms/step - loss: 0.2552 - val_loss:
0.2638
Epoch 72/100
9/9 _____ 0s 52ms/step - loss: 0.2577 - val_loss:
0.2657
Epoch 73/100
9/9 _____ 0s 24ms/step - loss: 0.2658 - val_loss:
0.2709
Epoch 74/100
9/9 _____ 0s 34ms/step - loss: 0.2678 - val_loss:
0.2618
```

Epoch 75/100
9/9 _____ 1s 47ms/step - loss: 0.2816 - val_loss: 0.2689
Epoch 76/100
9/9 _____ 1s 50ms/step - loss: 0.2666 - val_loss: 0.2618
Epoch 77/100
9/9 _____ 0s 22ms/step - loss: 0.2580 - val_loss: 0.2627
Epoch 78/100
9/9 _____ 0s 28ms/step - loss: 0.2622 - val_loss: 0.2600
Epoch 79/100
9/9 _____ 1s 22ms/step - loss: 0.2639 - val_loss: 0.2531
Epoch 80/100
9/9 _____ 0s 47ms/step - loss: 0.2613 - val_loss: 0.2548
Epoch 81/100
9/9 _____ 0s 28ms/step - loss: 0.2704 - val_loss: 0.2652
Epoch 82/100
9/9 _____ 0s 26ms/step - loss: 0.2516 - val_loss: 0.2486
Epoch 83/100
9/9 _____ 0s 28ms/step - loss: 0.2519 - val_loss: 0.2556
Epoch 84/100
9/9 _____ 0s 22ms/step - loss: 0.2428 - val_loss: 0.2513
Epoch 85/100
9/9 _____ 0s 23ms/step - loss: 0.2462 - val_loss: 0.2573
Epoch 86/100
9/9 _____ 0s 37ms/step - loss: 0.2390 - val_loss: 0.2446
Epoch 87/100
9/9 _____ 0s 37ms/step - loss: 0.2650 - val_loss: 0.2569
Epoch 88/100
9/9 _____ 0s 22ms/step - loss: 0.2331 - val_loss: 0.2543
Epoch 89/100
9/9 _____ 1s 62ms/step - loss: 0.2603 - val_loss: 0.2603
Epoch 90/100
9/9 _____ 0s 32ms/step - loss: 0.2608 - val_loss: 0.2493
Epoch 91/100

```

9/9 _____ 0s 29ms/step - loss: 0.2414 - val_loss:
0.2531
Epoch 92/100
9/9 _____ 0s 25ms/step - loss: 0.2463 - val_loss:
0.2587
Epoch 93/100
9/9 _____ 0s 23ms/step - loss: 0.2702 - val_loss:
0.2448
Epoch 94/100
9/9 _____ 0s 13ms/step - loss: 0.2745 - val_loss:
0.2446
Epoch 95/100
9/9 _____ 0s 13ms/step - loss: 0.2586 - val_loss:
0.2476
Epoch 96/100
9/9 _____ 0s 12ms/step - loss: 0.2462 - val_loss:
0.2345
Epoch 97/100
9/9 _____ 0s 13ms/step - loss: 0.2543 - val_loss:
0.2422
Epoch 98/100
9/9 _____ 0s 12ms/step - loss: 0.2648 - val_loss:
0.2376
Epoch 99/100
9/9 _____ 0s 15ms/step - loss: 0.2475 - val_loss:
0.2461
Epoch 100/100
9/9 _____ 0s 12ms/step - loss: 0.2538 - val_loss:
0.2970

X_test_reconstructed = autoencoder.predict(X_test)
reconstruction_errors = np.mean(np.square(X_test -
X_test_reconstructed), axis=1)
X_train_reconstructed = autoencoder.predict(X_train)
train_errors = np.mean(np.square(X_train - X_train_reconstructed),
axis=1)

141/141 _____ 0s 2ms/step
9/9 _____ 0s 4ms/step

thresholds = [85, 90, 95]
for t in thresholds:
    threshold = np.percentile(train_errors, t)
    y_pred = (reconstruction_errors > threshold).astype(int)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

```

```

print(f"\n Threshold: {t}th percentile")
print(f"Threshold value: {threshold:.4f}")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f} | Recall: {rec:.4f} | F1: {f1:.4f}")

```

```

 Threshold: 85th percentile
Threshold value: 0.3146
Accuracy : 0.8822
Precision: 0.7799 | Recall: 0.9989 | F1: 0.8759

```

```

 Threshold: 90th percentile
Threshold value: 0.3730
Accuracy : 0.9136
Precision: 0.8289 | Recall: 0.9984 | F1: 0.9058

```

```

 Threshold: 95th percentile
Threshold value: 0.5298
Accuracy : 0.9522
Precision: 0.8993 | Recall: 0.9968 | F1: 0.9456

```

Scores per threshold

```

thresholds = ['85%', '90%', '95%']
precisions = [0.7799, 0.8289, 0.8993]
recalls = [0.9989, 0.9984, 0.9968]
f1_scores = [0.8759, 0.9058, 0.9456]

```

Reorganize for metric-based grouping

```

metrics = ['Precision', 'Recall', 'F1-Score']
scores = [precisions, recalls, f1_scores]

```

```

x = np.arange(len(metrics))
width = 0.25

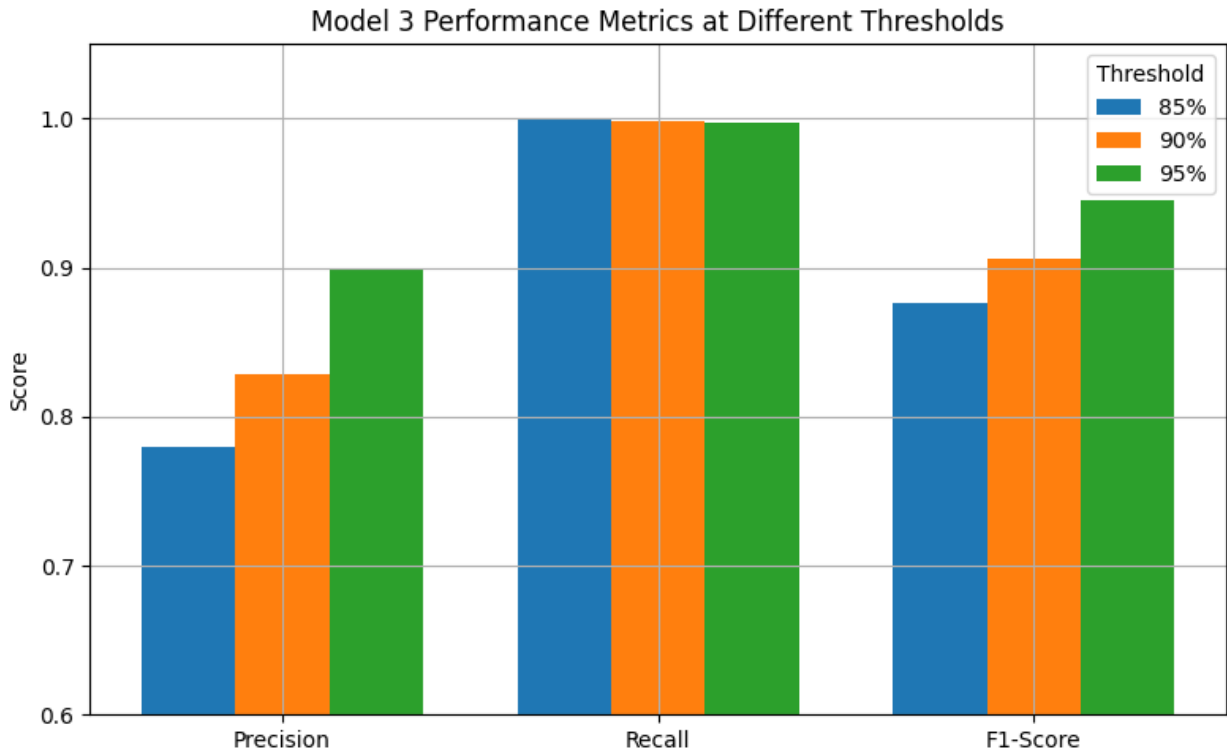
```

```

plt.figure(figsize=(8, 5))
plt.bar(x - width, [s[0] for s in scores], width, label='85%')
plt.bar(x, [s[1] for s in scores], width, label='90%')
plt.bar(x + width, [s[2] for s in scores], width, label='95%')

plt.xticks(x, metrics)
plt.ylim(0.6, 1.05)
plt.ylabel('Score')
plt.title('Model 3 Performance Metrics at Different Thresholds')
plt.legend(title='Threshold')
plt.grid(True)
plt.tight_layout()
plt.show()

```

Summary:

Project Summary:

This project applies unsupervised deep learning techniques to detect anomalies in ECG signals from the ECG5000 dataset. The approach is based on autoencoders, which are trained to learn and reconstruct patterns of normal heartbeats. The reconstruction error is used to identify anomalous signals that deviate from this learned normal pattern.

Models Developed:

1. Model 1 (Baseline): Trained on both normal and abnormal data.
 - Weak anomaly separation due to learning to reconstruct everything.
 2. Model 2 (Clean): Trained only on class '1' (normal data).
 - Significantly improved performance by focusing on normal patterns only.
 3. Model 3 (Optimized): Trained on normal data with Dropout and L2 regularization.
 - Achieved the best results, with high precision, recall, and F1-score.
-

Evaluation Metrics:

- Reconstruction error used to identify anomalies.
- Thresholds (85%, 90%, 95%) tested to control sensitivity.

- Metrics: Precision, Recall, F1-Score, Accuracy, ROC-AUC.
-

Why Normal-Only Training Works:

By training exclusively on normal data, the model becomes highly sensitive to anything unfamiliar (anomalies), resulting in higher reconstruction error. This improves anomaly detection accuracy and ensures better separation between normal and abnormal signals.

Conclusion:

Autoencoders are highly effective for anomaly detection when trained only on clean data. Proper threshold tuning and regularization significantly improve detection performance.